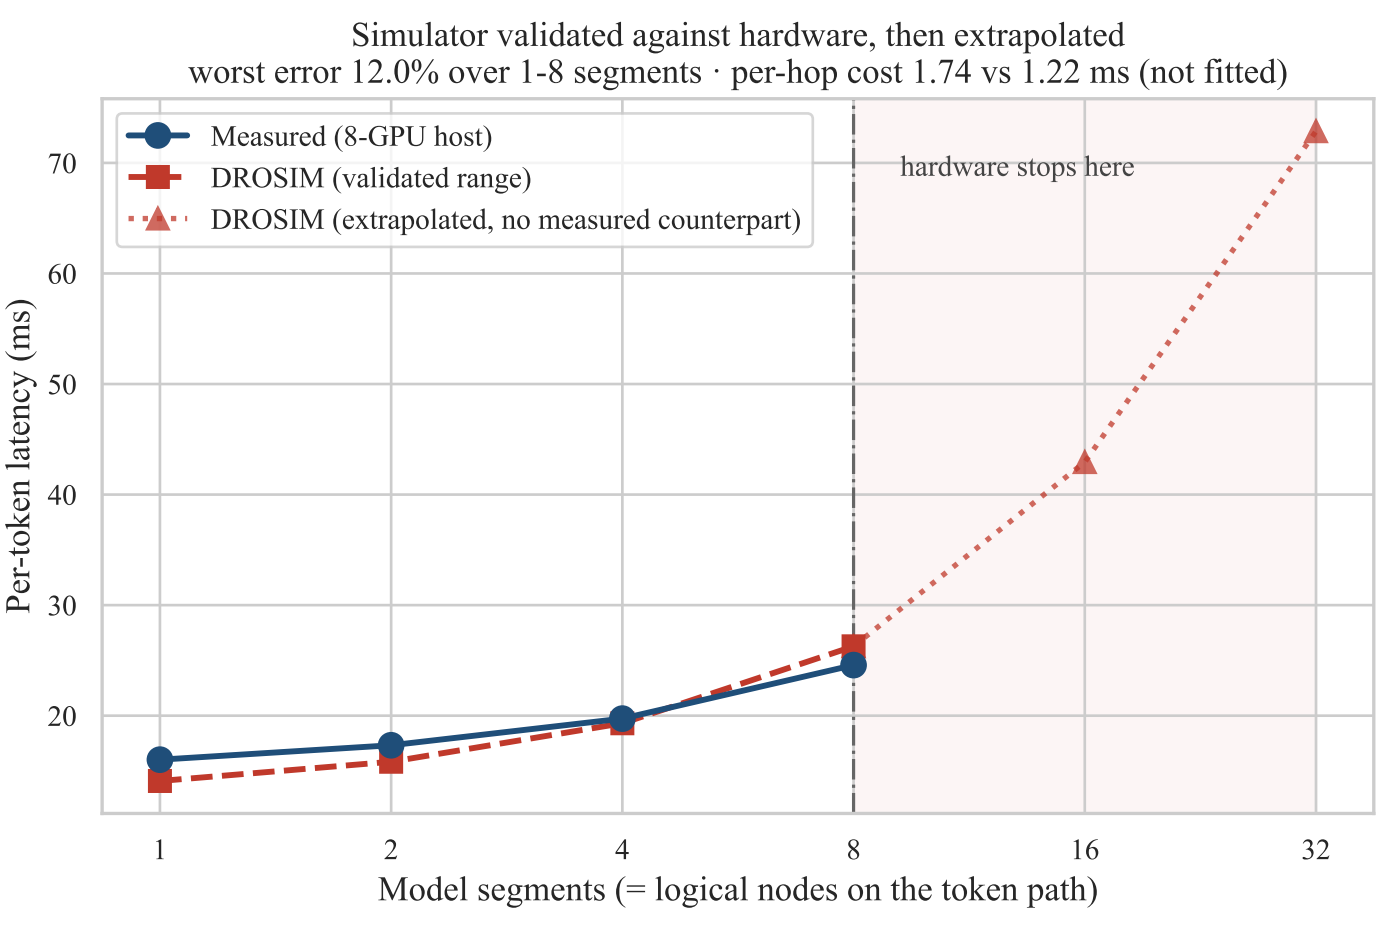
Simulator validated against hardware, then extrapolated
worst error 12.0% over 1-8 segments · per-hop cost 1.74 vs 1.22 ms (not fitted)
Measured (8-GPU host)
DROSIM (validated range)
DROSIM (extrapolated, no measured counterpart)
hardware stops here
Per-token latency (ms)
70
60
50
40
30
20
1
2
4
8
16
32
Model segments (= logical nodes on the token path)